

level 2 data collections in zarr

CNG Forum 2026 · Thu, Oct 08, 2026



Tom Nicholas · Earthmover



Sean Harkins · Development Seed



OUTLINE

what we'll cover


1. Level 2 vs Level 3 today
2. Features of an ideal cloud-native data system
3. Limitations of COG + STAC
4. Idea: Icechunk DataCollections
5. (Sean) Prototypes and future work



CNG FORUM 2025


last year: zarr vs cog

Lots of debate last year, and the rough consensus was:



LEVEL 3 · GRIDDED DATACUBES

Zarr wins



LEVEL 2 · INDIVIDUAL SCENES

COG + STAC wins

but similar aims: so why the split?



DESIRED FEATURES

icechunk-zarr works great! (for level 3)

Property	Icechunk-Zarr
Single endpoint  	✓
Query by coordinates  	✓
Scalable  	✓
Serverless  	✓
Uncoordinated reads  	✓
Uncoordinated writes 	✓

 i.e. "cloud-optimized"

 from Zarr

Property	Icechunk-Zarr
Arbitrary N-D schemas 	✓
Domain-agnostic 	✓
ACID transactions and consistency 	✓
Versioning, time travel, branches 	✓
Schema evolution 	✓
Zero-copy ingestion ("virtual chunks") 	✓

 from Xarray

 from Icechunk



LEVEL 2 VS LEVEL 3



but level 2 in zarr is unsolved

- You **cannot** think of Level 2 data as one datacube
- Problem: **many arrays, no shared coordinate system**
- So geospatial folks use many separate COGs, then index them with STAC
 - We also see anti-patterns: many small Zarrs, or one Zarr with too many groups

Important

This problem is **not** specific to geospatial 🌍 !

- 🔬 Bioimaging: millions of OME-TIFFs + an index.
- 🔭 Astronomy: many images as separate FITS files.
- 🔗 Fusion: many Zarr stores + a Parquet index.



different footprints and CRSs: nothing shared to stack along



FEATURE COMPARISON

stac + cog solves some problems, but not all

Property	Icechunk-Zarr	STAC + COG
Single endpoint 	✓	✓
Query by coordinates 	✓	✓
Scalable 	✓	✓
Serverless 	✓	● ¹
Uncoordinated reads 	✓	✓
Uncoordinated writes	✓	● ²

 i.e. "cloud-optimized"

Property	Icechunk-Zarr	STAC + COG
Arbitrary N-D schemas	✓	✗
Domain-agnostic	✓	✗
ACID transactions and consistency	✓	✗
Versioning, time travel, branches	✓	✗
Schema evolution	✓	✗
Zero-copy ingestion	✓	● ³

- ¹ only with stac-geoparquet (no STAC API server)
- ² only if writing to separate COG files
- ³ only if data already Cloud-Optimized (e.g. COG/Zarr)



IN PRACTICE

cog's limitations cause real problems

Consistency failures

During operational updates, the catalog and the data can disagree

Rigid data layouts

No hyperspectral bands, no ensemble members, no time dimension

Query-optimized chunking

Can't always rechunk for access patterns, e.g. timeseries

Expensive modifications

Renaming one variable in a 1TB TIFF means downloading and re-uploading 1TB

Domain-specific

Microscopists have near-identical requirements, but COG + STAC are geo-only


Existing archives

Can't reference non-cloud-optimized data, e.g. non-COG TIFFs



METADATA SEPARATE FROM DATA

disjoint systems: no consistency guarantees!



Data and metadata live in **two systems**, linked only by `href`s.

① Producer writes a new COG...

✨ ...but the job crashes before updating STAC → **orphaned data**, invisible to search

② Scene B is reprocessed and the old COG deleted → **dangling link**

Nothing ties the two together transactionally → **consistency problems**



SINGLE FILE FORMAT = DEAD END

these limitations are inherent to cog

A single-file container format can never support key features in object storage, since object storage has no edit-in-place.

- Prevents truly uncoordinated writes
- Prevents cheap updates → no schema evolution, no versioning
- **Prevents consistency(!)** using the serverless mechanism
Icechunk uses (conditional puts)
- Virtual references inherently involve multiple files

TIFF as a container constrains the schema


GeoTIFF + STAC are domain-specific

 rename one variable...

COG one 1 TB object



Icechunk-Zarr many small objects





IDEA: ICECHUNK DATACOLLECTIONS

better: many images + catalog in one repo

- What if we put **many images** (any shapes, any CRSs) into one Icechunk repo...
- ...alongside a searchable **STAC-like metadata table** that indexes them?
- Chunks and catalog are versioned together: **ACID, consistency, versioning and virtual chunks, for free**
- And it's domain-agnostic and works for any set of related arrays



one Icechunk repo



metadata catalog

id	datetime	EPSG
A	2024-07-02	32610
B	2024-07-03	32611
C	2024-07-05	32633
D	2024-07-09	3031




chunk storage

		A
		B
		C
		D

✓ versioned together · updated in one atomic commit

INITIAL ATTEMPT

Most of the "metadata" we're discussing can be modeled as an Arrow schema. So as an initial experiment we thought **What if we could store Arrow like data in Icechunk?**




Zarr-Datafusion-Search

github.com/developmentseed/zarr-datafusion-search

ZARR-DATAFUSION-SEARCH

Datafusion let's us build queryable database-like systems from any backend that can emit Arrow RecordBatches. With this we can represent columnar formats like Parquet using a series of 1-D Zarr arrays.



ZARR-DATAFUSION-SEARCH

Pros and cons

Pros

Only requires **Icechunk** and a compliant **DataFusion table provider**.

Materialized **R-tree** and **B-tree indexes** can be stored as Zarr arrays and used in the query pipeline, which can be much more efficient than common Parquet engine pushdown optimizations.


Writing data only requires an **Icechunk-compatible Zarr client**.

Cons

A very **custom solution** that doesn't leverage other great industry tools.

Writers need to coordinate writing their metadata "columns" so that "rows" are **chunk aligned**.

Variable-length dtypes have poor decoding performance in Zarr.



ICECHEST

The best of both worlds?

What if we could take advantage of the best parts of an **array store** like Icechunk and an **open table format** like Iceberg?



Icechunk

Array data

+



Apache Iceberg

Tabular metadata

=



icechest

arrays + tables


In one Icechunk commit

github.com/developmentseed/icechest

ICECHEST

How it works


Iceberg normally relies on a **catalog service** to track the current `metadata.json`. Instead of a catalog service, Icechest stores that pointer in **Icechunk commit metadata**.



ICECHEST PROS

Standard tooling, one commit



Rather than reinventing the wheel like we did with Zarr-Datafusion-Search, Icechest lets us use **standard tooling** for writing and reading both **columnar data** and **array data**.



ICECHEST CONS

Table maintenance is on us

A catalog service normally handles table maintenance. In Icechest, Iceberg versions are pinned by **Icechunk snapshots**, which have their **own lifecycle**, so we have to manage maintenance ourselves.



ICECHEST CONS

Icechunk node scaling

Level 3 data may have billions of chunks, but in **one array**. Level 2 data means **many discrete groups and arrays**. Virtualizing the full HLS archive into Icechest takes a mind-boggling number of **Zarr nodes**.

LEVEL 3



1 array node
billions of chunks

LEVEL 2 (HLS)

```

/
├── HLS.S30.T10SEG.2024001/
│   ├── B01/ multiscales
│   │   ├── 0
│   │   ├── 1
│   │   └── 2
│   ├── B02/ ...
│   └── Fmask/ ...
└── HLS.L30.T10SEG.2024003/ ...
    ... 38 M granules
    
```

181–217 nodes per granule
group/ array

collection	granules	nodes/granule	total nodes
HLSL30	16,102,070	181	2.91 B
HLSS30	22,051,832	217	4.79 B
combined	38,153,902		7.70 B

Commit memory and time grow with **total nodes in the store** ([icechunk#2449](#)): one year of HLS (~318 M nodes) → **~499 GiB, ~5.7 min** per commit.

This is fixable: Icechunk's **manifest splitting** approach just needs to be applied to the **node map**, and we're working on it now.

level 2 Icechunk data collections

what do you think?

 slides.earthmover.io/main/cng-forum-2026-level2-datacollections



Tom Nicholas · Earthmover



Sean Harkins · Development Seed



BONUS

common objections




SINGLE-FILE DOWNLOAD

"but i like downloading a single file"

sure — but that needn't dictate the storage layout

- It's reasonable to want a file on your local filesystem
- But we're free to transform the data **any way we like** on its journey from the datacentre to your laptop
- So use a service (e.g. Flux, via OGC API – EDR) to subset the store and assemble a GeoTIFF on request
- A file that never existed on disk — **the perfect COG for your immediate need**






COMPATIBILITY

"but many applications understand tiff!"

yes — *filesystem* applications

- We can **reconstruct GeoTIFFs on demand** for any local program (e.g. ArcGIS)
- It's genuinely impressive that macOS Preview can open a COG — a feat of format stability
 - But no one points Preview at object storage
- Any application that reads TIFFs **from object storage** can be taught to read Zarr (it's all just range requests)
- Compatibility via adapter layer is a **much easier problem** than consistency - vibe code it!





EXISTING COGS

"what if my data is already in cog?"

great — ingest it *virtually*, no copying

- Read just the **TIFF headers** (KBs per file) to find where every tile lives
- Store those byte ranges as **virtual chunk references** in an Icechunk repo
 - e.g. VirtualiZarr + DevSeed's `virtual-tiff` parser
- The COGs stay exactly where they are — **zero bytes of imagery copied**
- ...and you still get the catalog, versioning and transactions on top of your existing archive



↓ read headers only → byte ranges



Icechunk repo chunk manifest

chunk	url	offset	length
A/0.0	s3://.../scene_A.tif	16384	65536
A/0.1	s3://.../scene_A.tif	81920	65536
B/0.0	s3://.../scene_B.tif	16384	65536

✓ catalog + versioning + transactions, on references



ROLE OF COG

"then what role is left for cog?"

a great *delivery* format

- A COG is a fantastic thing to **hand to someone**: one self-describing file that works with almost every tool
- But the **system of record** needs consistency, versioning and a catalog
- So: **store** in Icechunk, **deliver** as COG — exported or generated on demand
- Storage format ≠ delivery format



Icechunk repo

system of record: versioned, transactional, catalogued



deliver



COG exports

for distribution &
archives



on-demand COG

subset via Flux / EDR



map tiles

for web maps